

Relational Databases, SQL

Ahmet Sacan

What you'll need

- ~~Firefox, SQLite plugin~~
- SqliteStudio
- Matlab: Database toolbox or a SQLite database driver
- Python: No additional packages needed

DB Definitions

- Relational database: collection of *tables* (also called *relations*)
- Database Management System (DBMS): a software that is used to create, access, and maintain a database; e.g., *sqlite*, *mysql*, *postgresql*, *oracle*, *sql server*.
- Table: Collection of rows (also called *tuples* or *records*).
- Each row in a table contains a set of columns (also called *fields* or *attributes*).
- Each column has a type:
 - Text: `VARCHAR(20)`
 - Integer: `INTEGER`
 - Floating-point: `FLOAT`, `DOUBLE`
 - Date/time `DATE`, `TIME`, `DATETIME`
- Primary key: provides a unique identifier for each row (optional).
- Schema: the structure of the database tables
 - The table name
 - The names and types of its columns
 - Various optional additional information (defaults, constraints, etc.)

SQL

- SQL = "Structured Query Language"
 - Non-procedural
 - Set-oriented
 - Relationally complete
 - Functionally incomplete
- Four main types of queries:
 - Insert, Delete, Select, Update

Create/drop table

- Syntax:
 - https://www.sqlite.org/lang_createtable.html
- Create a table for the students:

```
CREATE TABLE students (  
  id INTEGER PRIMARY KEY,  
  name VARCHAR(30)    UNIQUE,  
  birth DATE,  
  gpa FLOAT,  
  grad INTEGER  
);
```

- Drop table:
DROP TABLE students;

The following failed to have auto-incremented ids: CREATE TABLE students(id INTEGER AUTO_INCREMENT, ..., PRIMARY KEY (id));

Insert/Delete

- **Add rows to the students table:**

```
INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Anderson', '1987-10-22', 3.9, 2009);
INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Jones', '1990-4-16', 2.4, 2012);
INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Hernandez', '1989-8-12', 3.1, 2011);
INSERT INTO students(name, birth, gpa, grad)
    VALUES ('Chen', '1990-2-4', 3.2, 2011);
```

- **Delete row(s):**

```
DELETE FROM students;
```

Select Queries

- 3 main elements:
 - What you want
 - Where it is found
 - How you want it filtered
- Show entire contents of a table:

```
SELECT * FROM students;
```

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-04-16	2.4	2012
3	Hernandez	1989-08-12	3.1	2011
4	Chen	1990-02-04	3.2	2011

Select

- Show just a few columns from a table:

```
SELECT name, gpa FROM students;
```

name	gpa
Anderson	3.9
Jones	2.4
Hernandez	3.1
Chen	3.2

- Filtering: only get a subset of the rows:

```
SELECT name, gpa FROM students WHERE gpa > 3.0;
```

name	gpa
Anderson	3.9
Hernandez	3.1
Chen	3.2

Select

- **Sorting:**

```
SELECT gpa, name, grad FROM students WHERE gpa > 3.0 ORDER BY gpa DESC;
```

gpa	name	grad
3.9	Anderson	2009
3.2	Chen	2011
3.1	Hernandez	2011

- **Limiting: only get a certain number of rows:**

```
SELECT name, gpa FROM students LIMIT 0,2;
```

name	gpa
Anderson	3.9
Jones	2.4

Update / Delete

- Update:

```
UPDATE students  
    SET gpa = 2.6, grad = 2013  
    WHERE id = 2;
```

- Delete:

```
DELETE FROM students  
    WHERE id = 2;
```

Joins

- Join: a query that merges the contents of 2 or more tables, retrieves information from the merged results.
- Join example: many-to-one relationship
- Students have advisors; add a new table describing faculty.

id	name	title
1	Fujimura	assocprof
2	Bolosky	prof

Join example: many-to-one relationship

- Add new column `advisor_id` to the `students` table. This is a *foreign key*.

id	name	birth	gpa	grad	advisor_id
1	Anderson	1987-10-22	3.9	2009	2
2	Jones	1990-04-16	2.4	2012	1
3	Hernandez	1989-08-12	3.1	2011	1
4	Chen	1990-02-04	3.2	2011	1

- Perform the join query: Get the students who are advised by Fujimura

```
SELECT s.name, s.gpa FROM students s, advisors a
WHERE a.name = 'Fujimura' AND a.id = s.advisor_id;
```

name	gpa
Jones	2.4
Hernandez	3.1
Chen	3.2

id	name	title
1	Fujimura	assocprof
2	Bolosky	prof

Join example: many-to-many relationship

- Courses: students take many courses, courses have many students
- Add a new table describing courses:

id	number	name	quarter
1	CS142	Web stuff	Winter 2009
2	ART101	Finger painting	Fall 2008
3	ART101	Finger painting	Winter 2009
4	PE204	Mud wrestling	Winter 2009

- Create a "join table" `courses_students` describing which students took which courses.

course_id	student_id
1	1
3	1
4	1
1	2
2	2
1	3
2	4
4	4

Join example: many-to-many relationship

- Find all students who took a particular course ('ART101'):

```
SELECT s.name, c.quarter
  FROM students s, courses c, courses_students cs
 WHERE c.number = 'ART101'
       AND c.id = cs.course_id
       AND cs.student_id = s.id;
```

name	quarter
Jones	Fall 2008
Chen	Fall 2008
Anderson	Winter 2009

id	name	birth	gpa	grad
1	Anderson	1987-10-22	3.9	2009
2	Jones	1990-04-16	2.4	2012
3	Hernandez	1989-08-12	3.1	2011
4	Chen	1990-02-04	3.2	2011

id	number	name	quarter
1	CS142	Web stuff	Winter 2009
2	ART101	Finger painting	Fall 2008
3	ART101	Finger painting	Winter 2009
4	PE204	Mud wrestling	Winter 2009

course_id	student_id
1	1
3	1
4	1
1	2
2	2
1	3
2	4
4	4

Additional Database Features

- **Indexes:** used to speed up searches
- **Concurrency:** Allow more than one program use the database at the same time.
- **Transactions:** used to group operations together to provide predictable behavior even when there are concurrent operations on the database.
- **Locks:** Used to limit concurrent access.
- **Views:** a virtual table for results of a stored query
- **Procedures:** a set of sql statements stored in the database

Exercise

- GO database
 - <http://geneontology.org/page/lead-database-schema>
 - <http://geneontology.org/sites/default/files/public/diag-godb-er.jpg>
 - <http://www.berkeleybop.org/goose/>
- Retrieve the names of the species that are under the genus 'Drosophila'
- Retrieve the genus and species name of the organisms whose species name has a prefix 'mel' (use LIKE function).
- Retrieve the gene symbols (gene_product.symbol) of all Drosophila melanogaster (species.genus, species.species) genes that are annotated to the 'nucleus'.
 - Join path: term.id -> graph_path.term1_id / term2_id -> association.term_id / gene_product_id -> gene_product.id
 - Join path: gene_product.species_id -> species.id